

Department I - C Plus Plus

Modern and Lucid C++ Advanced for Professional Programmers

Week 2 – Move Semantics

Thomas Corbat / Felix Morgner
Rapperswil, 01.03.2022
FS2022




- **Topics:**

- Motivation for Moving Elements
- Rvalue References
- Value Categories
- Special Member Functions
- Copy Elision

- You can explain why move semantics exist
- You can apply the different types of references that are available in C++
- You know the value categories of C++
- You can determine the value category of an expression
- You can implement the special member functions that move objects
- You know what copy elision is and it is mandatory and when it is optional

Motivation for Move Semantics






Create Thing
Copy Thing



```
struct CopyableThing {  
    CopyableThing() {  
        std::cout << "Create Thing\n";  
    }  
    CopyableThing(CopyableThing const &) {  
        std::cout << "Copy Thing\n";  
    }  
};  
  
CopyableThing create() {  
    CopyableThing t{};  
    return t;  
}  
  
int main() {  
    CopyableThing created = create();  
}
```



Create Thing
Move Thing




```
struct MoveableThing {  
    MoveableThing() {  
        std::cout << "Create Thing\n";  
    }  
  
    MoveableThing(MoveableThing &&) {  
        std::cout << "Move Thing\n";  
    }  
};  
  
MoveableThing create() {  
    MoveableThing t{};  
    return t;  
}  
  
int main() {  
    MoveableThing created = create();  
}
```

- Sometimes it is desirable to avoid copying values around for performance reasons


```
struct Planet {  
    //gigantic type with data on heap  
};  
  
std::vector<Planet> createPlanetsOfSolarSystem() {  
    std::vector<Planet> planets{};  
    Planet earth{};  
    planets.push_back(earth);           //copying a whole planet takes time  
    planets.push_back(createMars());    //creates a mars temporary; copy that too?  
    return planets;                    //copy everything once more?!  
}
```

- Is it really necessary to copy all those objects around?

● Copying



● Moving




```
struct ContainerForBigObject {
    ContainerForBigObject()
        : resource{std::make_unique<BigObject>()} {}

    ContainerForBigObject(ContainerForBigObject const & other)
        : resource{std::make_unique<BigObject>(*other.resource)} {}

    ContainerForBigObject(ContainerForBigObject && other)
        : resource{std::move(other.resource)} {}

    ContainerForBigObject & operator=(ContainerForBigObject const & other) {
        resource = std::make_unique<BigObject>(*other.resource);
        return *this;
    }

    ContainerForBigObject & operator=(ContainerForBigObject && other) {
        std::swap(resource, other.resource);
        //resource = std::move(other.resource) is possible too
        return *this;
    }

private:
    std::unique_ptr<BigObject> resource;
};
```

Rvalue References



● lvalue References

- Binds to an lvalue
- Syntax: T &
- The original must exist as long as it is referred to!

```
void modify(T & t) {  
    //manipulate t  
}  
  
void lvalueRefExample() {  
    T t = 5;  
    modify(t);  
    T & ir = t;  
    //...  
}
```

● rvalue References

- Binds to an rvalue
- Syntax: T &&
- Can extend the life-time of a temporary

```
T createT();  
  
void consume(T && t) {  
    //manipulate t  
}  
  
void rvalueRefExample() {  
    consume(T{});  
    T && t = createT();  
    //...  
}
```

- **An lvalue Reference is an alias for a variable**

- Binds an lvalue
- Syntax: T &
- The original must exist as long as it is referred to!

- **Can be used as**

- Function parameter type (most useful: no copy and side-effect on argument possible)
- Member or local variable (barely useful)
- Return type (Must survive!)

```
void increment(int & i) {  
    ++i; // side-effect on argument  
}
```

- Beware of dangling references: undefined behavior!



- **References for rvalues**

- Binds only rvalues
- Syntax: <Type> &&


- **Argument is either a literal, a temporary object or an explicitly converted lvalue**

```
void consume(Food && food);

Food fryBurger();

void fastFood() {
    Food fries{"salty and greasy"};
    consume(fryBurger());           //call with rvalue
    consume(fries);                 //cannot pass lvalue to rvalue reference
    consume(std::move(fries));      //explicit conversion lvalue to xvalue
    Food && burger = fryBurger();    //life-extension of temporary
}
```

Value Categories



- **CPL**

- lvalue: expression on the left-hand side of an assignment (memory location)
- rvalue: expression on the right-hand side of an assignment (value)

- **C++**

- A little more complicated
- lvalue: has identity
- rvalue: does not have identity (temporaries and literals)

- **Example why lvalue does not always mean “can be on the left-hand side of an assignment”**

```
int a = 0;  
a = 7; //a is an lvalue  
      //ok
```

```
int const a = 0;  
a = 7; //a is still an lvalue  
      //not ok, const lvalue
```

- **Example why rvalue does not always mean “cannot be on the left-hand side of an assignment”**
 - Not useful, but valid. `S{}` clearly is a temporary

```
#include <iostream>
#include <string>

struct S {
    S & operator=(std::string const & s) {
        std::cout << "got \"" << s << "\" assigned\n";
        return *this;
    }
};

int main() {
    S{} = "new value";
}
```

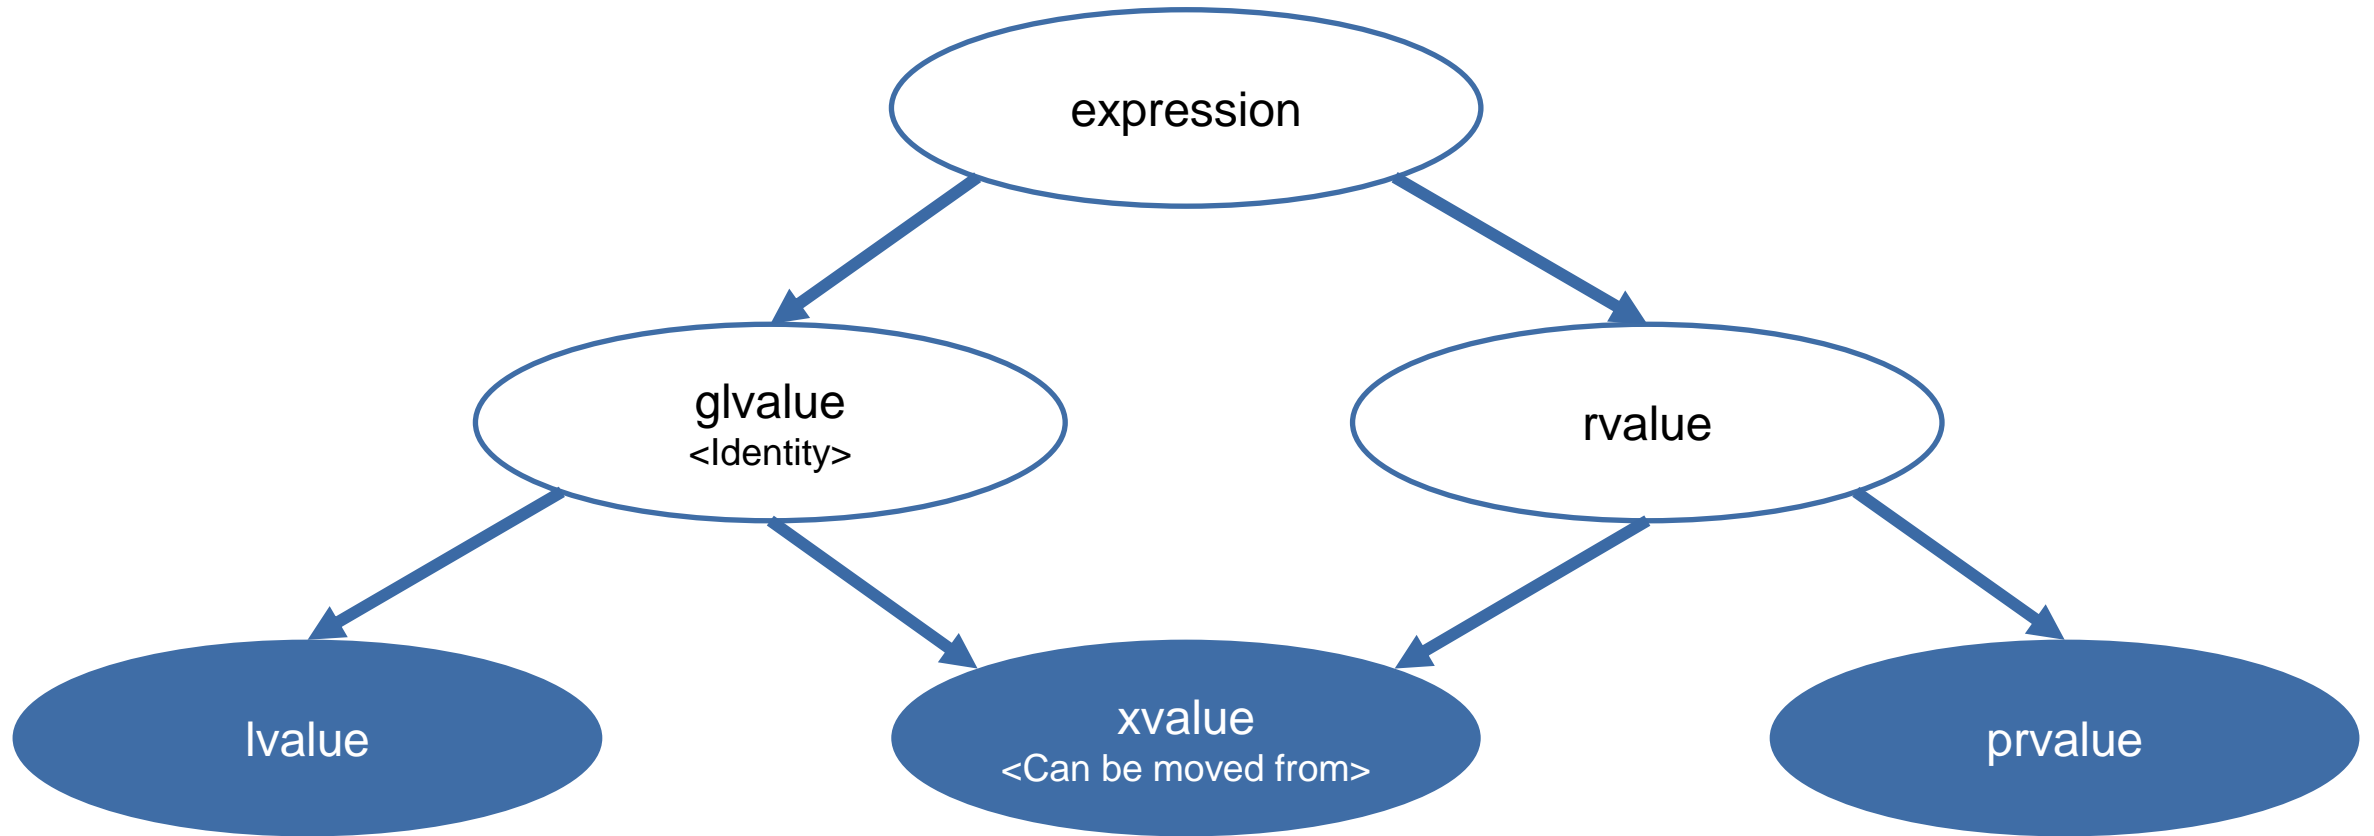

- **Every expression has**

- (non-reference) Type
- Value Category


- **Properties of a Value Category**

- has identity (address can be taken)
- can be moved from


| has identity? | can be moved from? | Value Category |
|---------------|--------------------|---------------------------|
| Yes | No | lvalue |
| Yes | Yes | xvalue (expiring value) |
| No | No (Since C++17) | prvalue (pure rvalue) |
| No | Yes (Since C++17) | - (doesn't exist anymore) |



- **Address can be taken**
- **Can be on the left-hand side of an assignment if modifiable (i.e. non-const)**
- **Can be used to initialize an lvalue reference**
- **Examples**
 - Names of variables and parameters (counter)
 - Function call with return type of lvalue reference to class type (`std::cout << 23`)
 - Built-in prefix increment/decrement expressions (`++a`)
 - Array index access (`arr[0]`), wenn `arr` is an lvalue
 - All string-literals by definition ("`name`")
 - This does not include user-defined (string) literals, like "`name`"s or "`name`"sv




- **Name: pure rvalue**
- **Address cannot be taken**
- **Cannot be left-hand side argument of built-in assignment operators**
- **Temporary materialization when a glvalue is required**
 - Conversion to xvalue
- **Examples:**
 - Literals: 23, false, nullptr, ...
 - Function call expression of non-reference return type: `int std::abs(int n)`
 - Several operators for built-in types, like post-increment/-decrement expressions: `x++`



- **Getting from something imaginary to something you can point to**
- **Prvalue to xvalue conversion happens...**
 - ... when binding a reference to a prvalue ①
 - ... when accessing a member of a prvalue ②
 - ... when accessing an element of a prvalue array
 - ... when converting a prvalue array to a pointer
 - ... when initializing an `std::initializer_list<T>` from a braced-init-list
- **Requires type to be complete and have a destructor**

```
struct Ghost {  
    void haunt() const {  
        std::cout << "booooo!\n";  
    }  
    //~Ghost() = delete;  
};  
  
Ghost evoke() {  
    return Ghost{};  
}  
  
int main() {  
    Ghost && sam = evoke(); ①  
  
    Ghost{}.haunt(); ②  
}
```


- **Name: expiring value**
- **Address cannot be taken**
- **Cannot be used as left-hand operator of built-in assignment**
- **Conversion from prvalue through temporary materialization**
- **Examples:**
 - Function call with rvalue reference return type, like `std::move`: `std::move(x)`
 - Access of non-reference members of an rvalue object
 - Array index access (`arr[0]`), wenn `arr` is an rvalue



```
X x1{}, x2{};
consume(std::move(x1));
std::move(x2).member;
X{}.member;
```

- **An lvalue Reference is an alias for a variable**

- Syntax: T &
- The original must exist as long as it is referred to!



- **Can be used as**

- Function parameter type (most useful: no copy and side-effect on argument possible)
- Member or local variable (barely useful)
- Return type (Must survive!)


```
void increment(int & i) {  
    ++i; // side-effect on argument  
}
```

- Beware of dangling references: undefined behavior!



- **References for rvalues**

- Syntax: T &&
- Binds to an rvalue (xvalue or prvalue)



- **Argument is either a literal or a temporary object**

```
std::string createGlass();

void fancy_name_for_function() {
    std::string mug{"cup of coffee"};
    std::string && glass_ref = createGlass(); //life-extension of temporary
    std::string && mug_ref = std::move(mug);  //explicit conversion lvalue to rvalue
    int && i_ref = 5;                        //binding rvalue reference to prvalue
}
```

- **Beware: Parameters and variables declared as rvalue references are lvalues in the context of function bodies! (Everything with a name is an lvalue)**
- **Beware 2.0: T&&/auto&& is not always an rvalue reference! (We'll come to that later)**


```
T value{};
std::cout << value;
```

```
int value{};
std::cout << value + 1;
```

```
void foo(T & param) {
    std::cout << param;
}
```

```
void print(T && param) {
    std::cout << param;
}
```

```
T create();
create();
```

```
T & create();
create();
```

```
T && create();
create();
```

```
T value{};
std::cout << value + 1;
```

```
T value{};
T o = std::move(value);
```


```
std::cout << "Hello";
```

| | |
|--|--------|
| <pre>T value{}; std::cout << <u>value</u>;</pre> | lvalue |
| <pre>int value{}; std::cout << <u>value + 1</u>;</pre> | rvalue |
| <pre>void foo(T & param) { std::cout << <u>param</u>; }</pre> | lvalue |
| <pre>void print(T && param) { std::cout << <u>param</u>; }</pre> | lvalue |
| <pre>T create(); <u>create()</u>;</pre> | rvalue |

| | |
|--|-----------------|
| <pre>T & create(); <u>create()</u>;</pre> | lvalue |
| <pre>T && create(); <u>create()</u>;</pre> | rvalue |
| <pre>T value{}; std::cout << <u>value + 1</u>;</pre> | depends on + |
| <pre>T value{}; T o = <u>std::move(value)</u>;</pre> | rvalue |
| <pre>std::cout << <u>"Hello"</u>;</pre> | lvalue |

● lvalue Reference

■ binds




```
void f(Type &);
```

```
Type t{};
f(t);
```

● const lvalue Reference

■ binds




```
void f(Type const &);
```

```
Type t{};
f(t);
f(std::move(t));
f(Type{});
```

● rvalue Reference

■ binds



```
void f(Type &&);
```

```
Type t{};
f(Type{});
f(std::move(t));
```


| | <code>f(S)</code> | <code>f(S &)</code> | <code>f(S const &)</code> | <code>f(S &&)</code> |
|--|-------------------|----------------------------|-------------------------------|------------------------------|
| <code>S s{};</code> <code>f(s);</code> | ✓ | ✓ (preferred over const &) | ✓ | ✗ |
| <code>S const s{};</code> <code>f(s);</code> | ✓ | ✗ | ✓ | ✗ |
| <code>f(S{});</code> | ✓ | ✗ | ✓ | ✓ (preferred over const &) |
| <code>S s{};</code> <code>f(std::move(s));</code> | ✓ | ✗ | ✓ | ✓ (preferred over const &) |

- The overload for value parameters imposes ambiguities.
- For deciding between two lvalue reference overloads (const and non-const) the constness of the argument is considered.

| | <code>S::m()</code> | <code>S::m() const</code> | <code>S::m() &</code> | <code>S::m() const &</code> | <code>S::m() &&</code> |
|---|---------------------|---------------------------|----------------------------------|---------------------------------|----------------------------------|
| <code>S s{};</code> <code>s.m();</code> | ✓ | ✓ | ✓ (preferred over const &) | ✓ | ✗ |
| <code>S const s{};</code> <code>s.m();</code> | ✗ | ✓ | ✗ | ✓ | ✗ |
| <code>S{}.m();</code> | ✓ | ✓ | ✗ | ✓ | ✓ (preferred over const &) |
| <code>S s{};</code> <code>std::move(s).m();</code> | ✓ | ✓ | ✗ | ✓ | ✓ (preferred over const &) |

- Reference and non-reference overloads cannot be mixed!
- The reference qualifier affects the this object and the overload resolution
- `const &&` would theoretically be possible, but it is an artificial case

Special Member Functions



- **Constructors**

- Default Constructor
- Copy Constructor
- Move Constructor

- **Assignment Operators**

- Copy Assignment
- Move Assignment

- **Destructor**

- **Advice:** If possible, design your types in way that the default implementations work for them. Library developers might need to implement custom special member functions.

```
struct S {  
    S();  
    ~S();  
    S(S const &);  
    S & operator=(S const &);  
    S(S &&);  
    S & operator=(S &&);  
};
```

● Move Constructor (Since C++11)

- Takes the entrails out of the argument and moves them to the constructed object
 - Leaves argument in valid but indeterminate state
 - Don't use the argument after it has been moved from until you assign it a new value
- Signature: && parameter of the same type
- Default Behavior (implicit or =default)
 - Initializes base-classes and members with move-initialization


S(S &&)




```
struct S {  
    S(S && s) : member{std::move(s.member)}  
    {...}  
    M member;  
};
```

```
void f(S param) {  
    S local{std::move(param)};  
    // don't use param until  
    // param = ...  
}
```


● Copy-Constructor



● Move-Constructor



● Copy Assignment Operator

- Copies the argument into the `this` object
- Assignment Operator with `const &` parameter of the same type
- Default Behavior (implicit or `=default`)
 - Initializes base-classes and members with copy-assignment

`S & operator=(S const &)`

```
struct S {  
    S & operator=(S const & s) {  
        member = s.member;  
        return *this;  
    }  
    M member;  
};
```

```
void f(S param) {  
    S local{};  
    local = param;  
    ...  
}
```

● Move Assignment Operator

S & operator=(S &&)

- Takes the entrails out of the argument and moves them to the this object
 - Leaves argument in valid but indeterminate state
 - Don't use the argument after it has been moved from until you assign it a new value
- Assignment Operator with && Parameter of the same Type
- Default Behavior (implicit or =default)
 - Assigns base-classes and members with move-assignment



```
struct S {  
    S & operator=(S && s) {  
        member = std::move(s.member);  
        return *this;  
    }  
    M member;  
};
```

```
void f(S param) {  
    S local{};  
    local = std::move(param);  
    ...  
}
```

● Destructor

~S()

- Deallocates resources held by the this object
- Signature: ~<Class-Name>()
- No Parameters
- Default Behavior (implicit or =default)
 - Calls destructor of base-classes and members
- Must not throw exceptions! (is noexcept)

```
struct S {  
    ~S() noexcept {...}  
    M member;  
};
```

<usually, you will not call
destructors explicitly>
Happens at end of scope: }



- **Assignment operators must be member functions**
- **Move operations must not throw exceptions**
 - They shall not allocate new memory
 - Otherwise `std::swap` won't work reliably
 - More on the topic of exception guarantees later
- **Use the default implementation whenever possible**

```
struct S {  
    S & operator=(S && s) noexcept;  
    S(S && other) noexcept;  
};
```

```
struct S {  
    S() = default;  
    ~S() = default;  
    S(S const &) = default;  
    S & operator=(S const &) = default;  
    S(S &&) = default;  
    S & operator=(S &&) = default;  
};
```

| | | What you get | | | | | Where you want to be |
|---------------------|----------------------|--------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| | default constructor | destructor | copy constructor | copy assignment | move constructor | move assignment | |
| nothing | defaulted | defaulted | defaulted | defaulted | defaulted | defaulted | |
| any constructor | not declared | defaulted | defaulted | defaulted | defaulted | defaulted | |
| default constructor | <u>user declared</u> | defaulted | defaulted | defaulted | defaulted | defaulted | |
| What you write | destructor | defaulted | <u>user declared</u> | defaulted (!) | defaulted (!) | not declared | not declared |
| | copy constructor | not declared | defaulted | <u>user declared</u> | defaulted (!) | not declared | not declared |
| | copy assignment | defaulted | defaulted | defaulted (!) | <u>user declared</u> | not declared | not declared |
| | move constructor | not declared | defaulted | deleted | deleted | <u>user declared</u> | not declared |
| | move assignment | defaulted | defaulted | deleted | deleted | not declared | <u>user declared</u> |
| | | | | | | | Avoid if possible |

Howard Hinnant's Table: https://accu.org/content/conf2014/Howard_Hinnant_Accu_2014.pdf

Note: Getting the defaulted special members denoted with a (!) is a bug in the standard.

- There are three different kinds of expression types in C++ (lvalue, xvalue, prvalue)
- The compiler must omit certain copy and move operations related to initialization from prvalues
- Objects/values can be copied, moved or passed by reference
- Good read about rvalue references and move semantics (state pre C++17):
http://thbecker.net/articles/rvalue_references/section_01.html
- Interesting talk about the problems with move semantics (by Nicolai Josuttis):
https://www.youtube.com/watch?v=PNRju6_yn3o

Implementing Assignment

Self-Study



- **The C++ Core Guidelines recommend providing a `swap()` member function and free function for value-like types (C.83)**
 - This enables the Copy-Swap-Idiom
 - `using std::swap` allows to fall back on the `std::swap` implementation if no user defined `swap()` is available for the member

```
struct S {  
    void swap(S & other) noexcept {  
        using std::swap;  
        swap(member1, other.member1);  
        //...  
    }  
    //...  
};  
  
void swap(S & lhs, S & rhs) noexcept {  
    lhs.swap(rhs);  
}
```

- **Usually you don't need to do this at all, but if you have to following this pattern is usually recommended**
 - Avoid self-copy
 - Use the copy constructor to create the copy of the argument
 - Swap the this-object with the copy (swapping is expected to be efficient)
 - Copy-Swap-Idiom

```
struct S {  
    S & operator=(S const & s) {  
        if (&s != this) {  
            S copy = s;  
            swap(copy);  
        }  
        return *this;  
    }  
    //...  
};
```

- Usually you don't need to do this at all
- If you have to, following this pattern is usually recommended
 - Avoid self-move
 - Swap the this-object with the parameter (swapping is expected to be efficient)

```
struct S {  
    S & operator=(S && s) {  
        if (&s != this) {  
            swap(s);  
        }  
        return *this;  
    }  
    //...  
};
```

Copy Elision

```
struct S {  
    S(S const & s) {  
        //Why is this not called?!  
    }  
};
```

Self-Study



- **In some cases the compiler is required to elide (omit) specific copy/move operations (regardless of the side-effects of the corresponding special member functions!)**
 - The omitted copy/move special member functions need not exist
 - If they exist, their side-effects are ignored
- **In initialization, when the initializer is a prvalue**
 - `S{}` is materialized in `s`
- **When a function call returns a prvalue (simplified)**
 - `S{}` is materialized in `new_sw`
 - `S{}` is materialized at the memory location return by `new`
We will cover explicit memory management later

```
S s = S{S{}};
```

```
S create() {  
    return S{};  
}  
  
int main() {  
    S new_sw{create()};  
    S * sp = new S{create()};  
}
```

- In some cases the compiler is allowed to further optimize specific copy/move operations (regardless of the side-effects of the corresponding special member functions!)

- Named return value optimization

```
S create() {  
    S s{};  
    return s;  
}  
  
int main() {  
    S s{create()};  
    s = create();  
}
```

- The constructors must still exist – even if they are elided.

```
int main() {
    std::cout << "\t --- S s{create()} ---\n";
    S s{create()};
    std::cout << "\t --- s = create() ---\n";
    s = create();
}
```

```
S create() {
    S s{};
    std::cout << "\t --- create() ---\n";
    return s;
}
```

Disabled elision (C++14):
-fno-elide-constructors

```
--- S s{create()} ---
Constructor S()
--- create() ---
Constructor S(S &&)
Constructor S(S &&)
--- s = create() ---
Constructor S()
--- create() ---
Constructor S(S &&)
operator =(S &&)
```

Disabled elision (C++17):
-fno-elide-constructors

```
--- S s{create()} ---
Constructor S()
--- create() ---
Constructor S(S &&)

--- s = create() ---
Constructor S()
--- create() ---
Constructor S(S &&)
operator =(S &&)
```

With elision (C++17):

```
--- S s{create()} ---
Constructor S()
--- create() ---

--- s = create() ---
Constructor S()
--- create() ---

operator =(S &&)
```

- In throw expressions (Since C++11)

```
try {  
    throw S{7};  
} catch (...) {  
}
```

- In catch clauses (Since C++11)

```
try {  
    throw S{7};  
} catch (S s) {  
}
```

- Beware: The compiler is allowed to change observable behavior with this optimization!
- To be sure to avoid copies still catch by const &

- Is the following a good idea?

```
S create() {  
    S s{};  
    return std::move(s);  
}
```

- While it sounds not that bad it prevents copy elision

```
S create() {  
    S s{};    //ctor  
    return s;  
}  
  
void foo() {  
    auto s = create();  
} //dtor
```

```
S create() {  
    S s{};    //ctor  
    return std::move(s); //move ctor  
} //dtor  
  
void foo() {  
    auto s = create();  
} //dtor
```

- **NRVO (Named Return Value Optimization)**

- Return type is value type
- Return expression is a local variable (more or less) of the return type
 - `const` is ignored for the type comparison
- The object is constructed in the location of the return value (instead of moved or copied)

- **throw Expression**

- Return expression is a local variable (more or less) from the innermost surrounding try block (if any)
- The object is constructed in the location where it would be moved or copied

- **catch Clause**

- If the caught type is the same as the object thrown, it access the object directly (as if caught by reference)
 - Must not change the observed behavior (except constructors/destructors)

Life-Time Extension

Self-Study



- Life-time of a temporary can be extended by "const lvalue reference" or "rvalue reference"
- Extended life-time ends at the end of the block

```
struct Demon { /*...*/ };
```

```
Demon summon() {  
    return Demon{};  
}
```

```
void countEyes(Demon const &) { /*...*/ }
```

```
int main() {
```

```
    summon();
```

```
    countEyes(Demon{});
```

```
    Demon const & flaaghun = summon();
```

```
    Demon && laznik = summon();
```

```
} //flaaghun and laznik die
```

//Demon dies at the end of the statement

//Demon lives long enough for count_eyes to finish

//Life-time can be extended by const &

// -> flaaghun lives until end of block

//Life-time can also be extended by &&

// -> laznik lives until end of block



Dangling
References

- Extension of life-time is not transitive

```
Demon const & bloodMagic() {
    Demon breknok{};
    return breknok;
} //When blood_magic ends, breknok dies and will stay dead. All access will be Undefined Behavior!

Demon const & animate(Demon const & demon) {
    /*...*/
    return demon;
}

int main() {
    Demon const & breknok = blood_magic();    //You cannot keep demon from blood_magic alive!
    // -> Access to breknok would be Undefined Behavior
    Demon const & knoorus = animate(Demon{}); //You cannot keep demon passed through animate alive!
    // -> Access to knoorus would be Undefined Behavior
}
```