

Department I - C Plus Plus

Modern and Lucid C++ Advanced for Professional Programmers

Week 10 – Networking and Threads

Prof. Peter Sommerlad / Thomas Corbat
Rapperswil, 25.04.2019
FS2019



Recap Week 9



- Create a tag type for the unit

```
struct Kph;  
struct Mph;  
struct Mps;
```

- Create a quantity type template for speed

```
template <typename Unit>  
struct Speed {  
    constexpr explicit Speed(double value)  
        : value{value}{};  
    constexpr explicit operator double() const {  
        return value;  
    }  
private:  
    double value;  
};
```

- Add a speedCast function

```
template<typename Target, typename Source>
constexpr Speed<Target> speedCast(Speed<Source> const & source) {
    return Speed<Target>{ ConversionTraits<Target, Source>::convert(source) };
}
```

- Create a ConversionTraits class template

```
template<typename Target, typename Source>
struct ConversionTraits {
    constexpr static Speed<Target> convert(Speed<Source> sourceValue) = delete;
};
```

```
namespace velocity::literals {  
  
constexpr inline Speed<Kph> operator"" _kph(unsigned long long value) {  
    return Speed<Kph>{safeToDouble(value)};  
}  
  
constexpr inline Speed<Kph> operator"" _kph(long double value) {  
    return Speed<Kph>{safeToDouble(value)};  
}  
  
//...  
}
```

- You know the basics of communication with sockets
- You can implement applications that communicate over a network
- You can implement network IO with synchronous and asynchronous operation
- You get familiar with the C++ thread API

Networking in C++ (With ASIO)



- **Sockets are an abstraction of endpoints for communication**

- **TCP Sockets**

- Reliable
- Stream-oriented
- Requires connection setup

- **UDP Sockets**

- Packets might get lost or arrive out of order
- Datagram-oriented (max 65k)
- Sending/receiving possible without a connection

Primitive	Meaning
Socket	Create a new communication end point
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data over the connection
Close	Release the connection

• Server

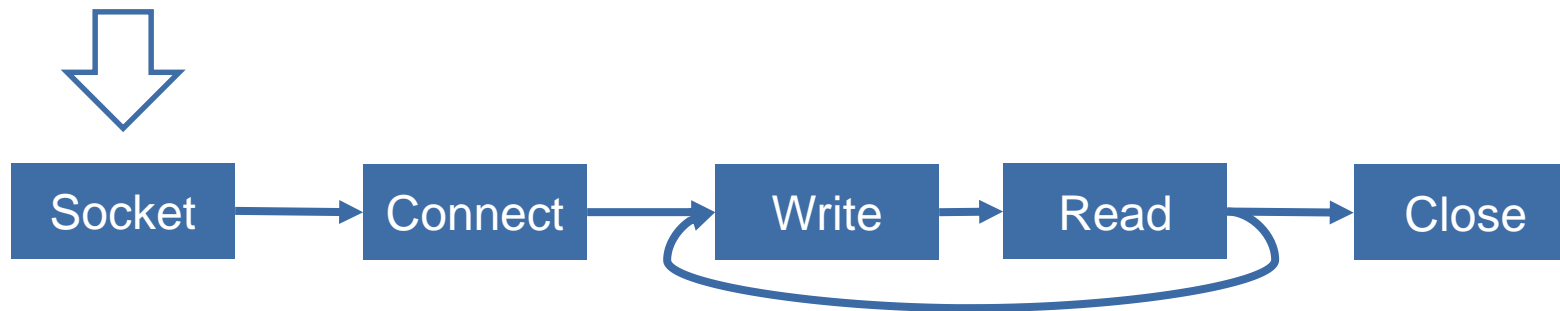


• Client



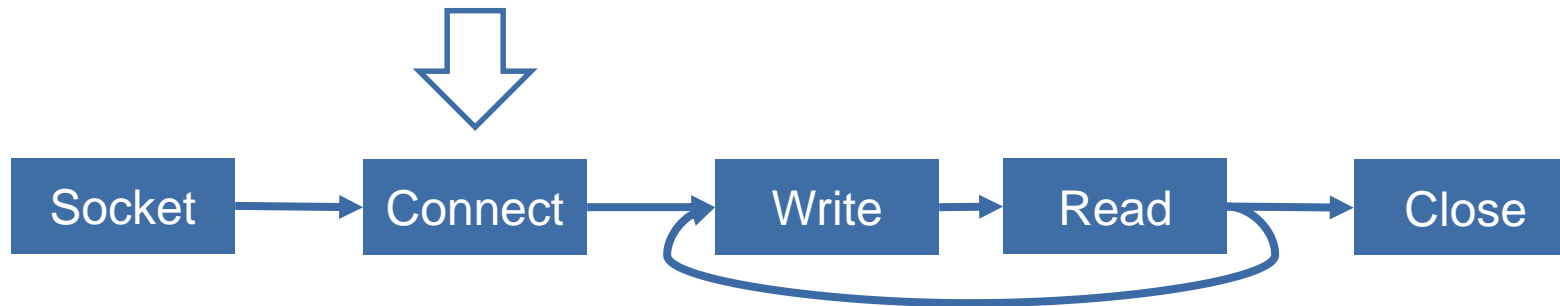
- All ASIO operations require an I/O context (more on that later)
- Create a TCP socket using the context

```
asio::io_context context{};  
asio::ip::tcp::socket socket{context};
```



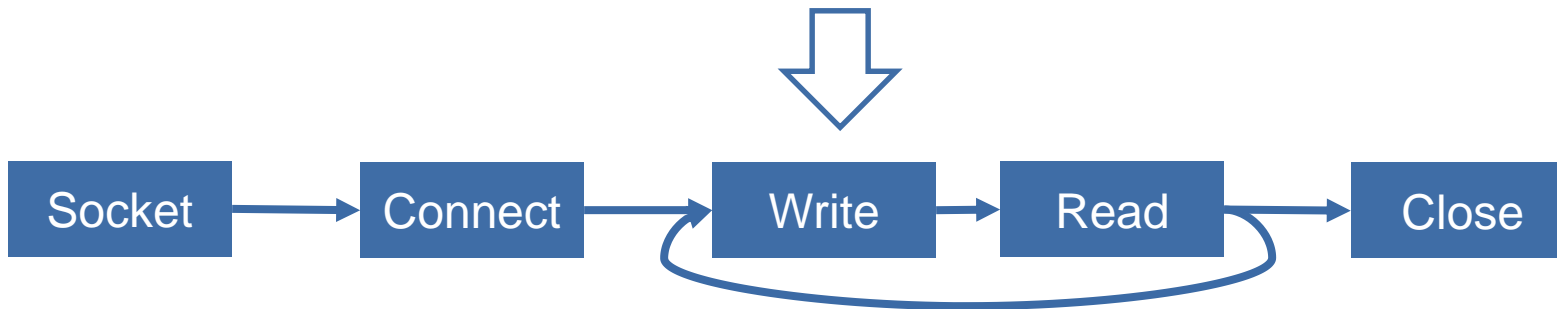
- A resolver resolves host names to end points (for TCP IP address and port)
- `asio::connect()` tries to establish a connection to given end point(s)
 - It tries `connect()` of the socket for each endpoint to establish a connection

```
asio::ip::tcp::resolver resolver{context};  
auto endpoints = resolver.resolve(domain, "80");  
  
asio::connect(socket, endpoints);
```



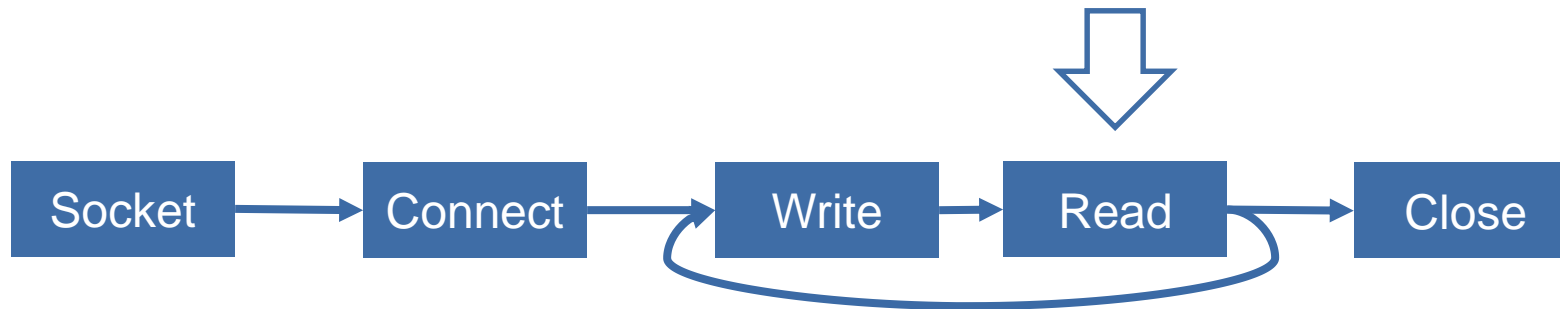
- **asio::write()** sends the data to the peer the socket is connected to
 - It returns when all data is sent or an error occurred (exception: asio::system_error)

```
std::ostream request{};  
request << "GET / HTTP/1.1\r\n";  
request << "Host: " << domain << "\r\n";  
request << "\r\n";  
  
asio::write(socket, asio::buffer(request.str()));
```



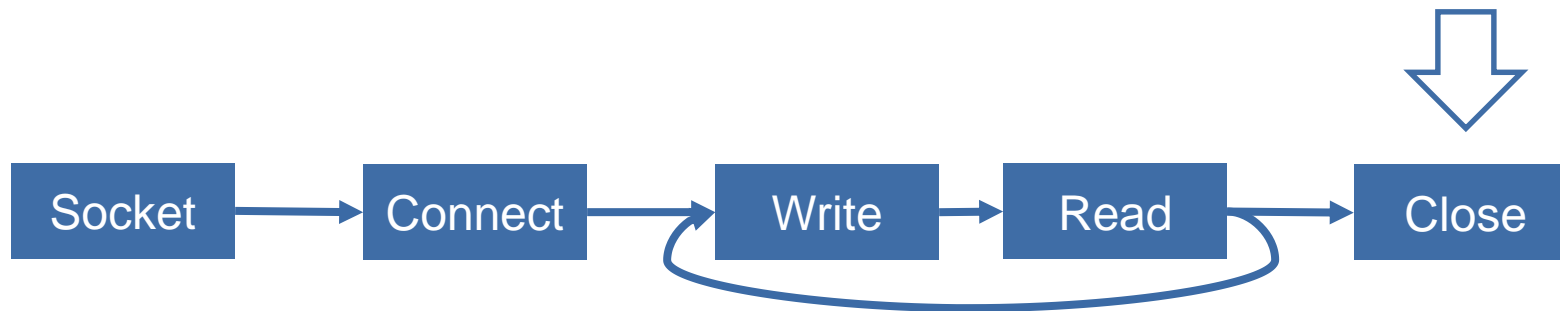
- **`asio::read()` receives data sent by the peer the socket is connected to**
 - It returns when all data is sent or an error occurred
- **The error code is set if a problem occurs or the stream has been closed (`asio::error::eof`)**

```
constexpr size_t bufferSize = 1024;  
std::array<char, bufferSize> reply{};  
asio::error_code errorCode{};  
auto readLength = asio::read(socket, asio::buffer(reply.data(), bufferSize), errorCode);
```




- `shutdown()` closes the read/write stream associated with the socket, indicating to the peer that no more data will be received/sent
- The destructor of the socket cancels all pending operations and destroys the object

```
socket.shutdown(asio::ip::tcp::socket::shutdown_both);  
socket.close();
```




- An acceptor is a special socket responsible for establishing incoming connections
- In ASIO the acceptor is bound to a given local end point and starts listening automatically

```
asio::io_context context{};  
asio::ip::tcp::endpoint localEndpoint{asio::ip::tcp::v4(), port};  
asio::ip::tcp::acceptor acceptor{context, localEndpoint};
```




- The `accept()` member function blocks until a client tries to establish a connection (with connect)
- It returns a new socket through which the connected client can be reached

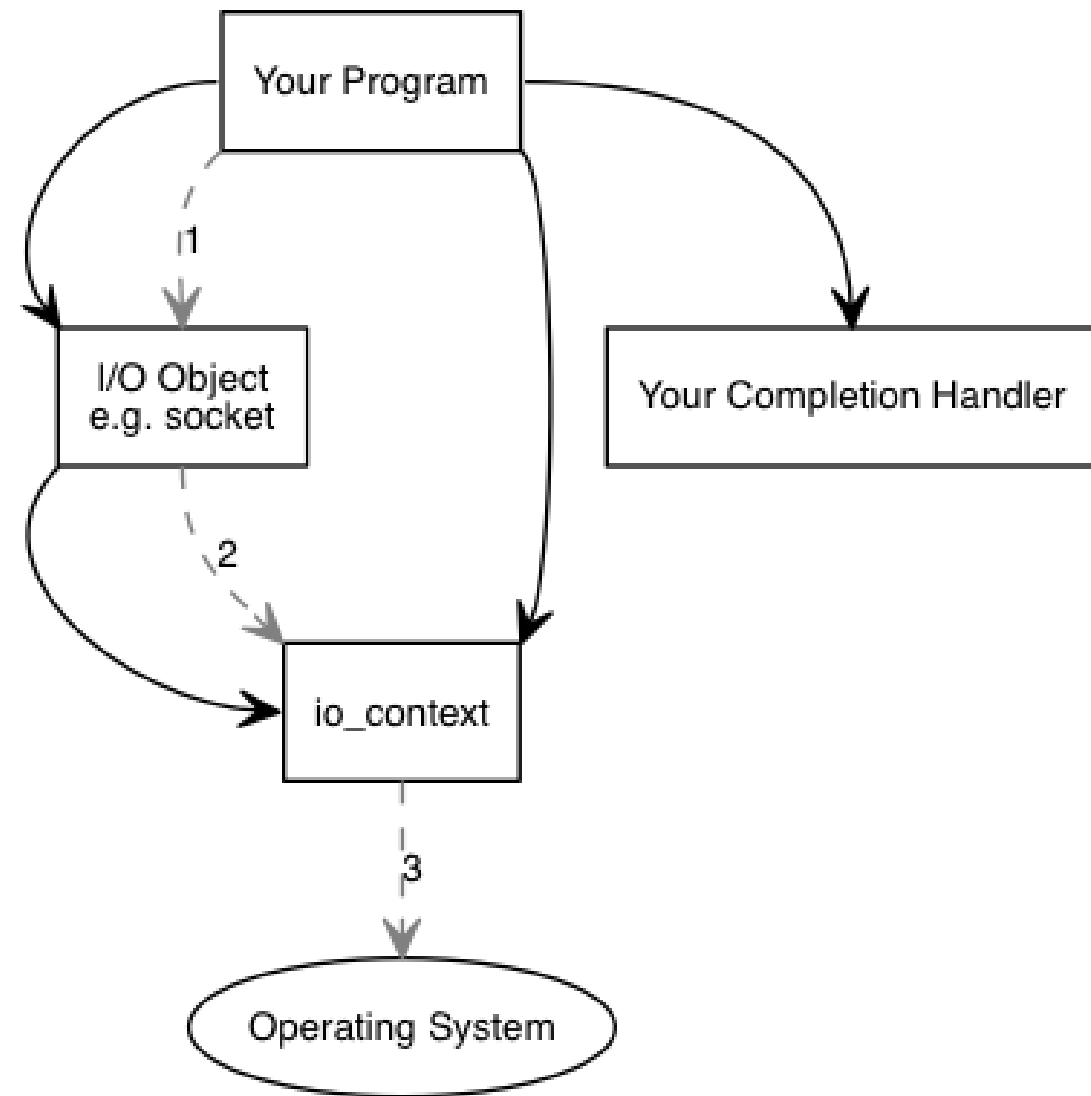
```
asio::ip::tcp::endpoint peerEndpoint{};  
asio::ip::tcp::socket peerSocket = acceptor.accept(peerEndpoint);
```




- **Using synchronous operations blocks the current thread**
- **No other operations can be executed while being blocked**
- **Asynchronous operations allow further processing of other requests while the async operation is executed (being blocked or making progress)**
- **Most operating systems support asynchronous IO operations**



1. The program invokes an async operation on an I/O object and passes a completion handler as a callback
2. The I/O object delegates the operation and the callback to its `io_context`
3. The operating system performs the asynchronous operation



4. The operating system signals the `io_context` that the operation has been completed
5. When the program calls `io_context::run()` the remaining asynchronous operations are performed (wait for the result of the operating system)
6. Still inside `io_context::run()` the completion handler is called to handle the result (or error) of the asynchronous operation



- **Async read operations**

- `asio::async_read`
- `asio::async_read_until`
- `asio::async_read_at`

- **Async write operations**

- `asio::async_write`
- `asio::async_write_at`

- **They return immediately**

- **The operations is processed by the executer associated with the stream's `asio::io_context`**

- **A completion handler is called when the operation is done**



- **asio::async_read_until** (the call returns immediately)
 - Reads from asynchronous stream
 - Into a buffer
 - Until a specific character is encountered
 - Then it calls the completion handler
- **The completion handler is a callable taking an asio::error_code and a std::size_t as arguments**

```
auto readCompletionHandler = [] (asio::error_code ec, std::size_t length) {  
    //...  
};  
  
asio::async_read_until(socket, buffer, '\n', readCompletionHandler);
```

- **asio::async_write (the call returns immediately)**
 - Writes to an asynchronous stream
 - The data from a buffer
 - Until all data has been written or an error occurs
 - Then it calls the completion handler
- **The completion handler is a callable taking an asio::error_code and a std::size_t as arguments**

```
auto writeCompletionHandler = [] (asio::error_code ec, std::size_t length) {  
    //...  
};  
  
asio::async_write(socket, buffer, writeCompletionHandler);
```

```
struct Server {  
    using tcp = asio::ip::tcp;  
    Server(asio::io_context & context, unsigned short port)  
        : acceptor{context, tcp::endpoint{tcp::v4(), port}}{  
        accept();  
    }  
  
private:  
    void accept() {  
        auto acceptHandler = [this] (asio::error_code ec, tcp::socket peer) {  
            if (!ec) {  
                auto session = std::make_shared<Session>(std::move(peer));  
                session->start();  
            }  
            accept();  
        };  
        acceptor.async_accept(acceptHandler);  
    }  
    tcp::acceptor acceptor;  
};
```



```
void accept() {  
    auto acceptHandler = [this] (asio::error_code ec, tcp::socket peer) {  
        if (!ec) {  
            auto session = std::make_shared<Session>(std::move(peer));  
            session->start();  
        }  
        accept();  
    };  
    acceptor.async_accept(acceptHandler);  
}
```

- **Creates an accept handler that is called when an incoming connection has been established**
 - The second parameter is the socket of the newly connected client
 - A Session object is created (on the heap) to handle all communication with the client
 - `accept()` is called to continue accepting new inbound connection attempts
- **The accept handler is registered to handle the next accept asynchronously**

```
Server(asio::io_context & context, unsigned short port)
    : acceptor{context, tcp::endpoint{tcp::v4(), port}}{
    accept();
}
```

- The constructor creates the server
- It initializes its acceptor with the given io_context and port
- It calls accept for registering the accept handler for the next incoming connection attempt
 - This function does not block

```
int main() {  
    asio::io_context context{};  
    Server server{context, 1234};  
    context.run();  
}
```

- **Create an io_context**
 - It has an associated executor that handles the asynchronous calls
- **Create the server on port 1234**
- **Run the executor of the io_context until no async operation is left**
 - Since we already have an `async_accept` request pending this operation does not return immediately
 - We will keep the this `run()` call busy
- **It is important that the server object lives as long as async operations on it are processed**

- **Constructor**

- Stores the socket with the client connection

- **start() initiates the first async read**

- **read() invokes async reading**

- **write() invokes async writing**

- Called by the handler in read

- **The fields store the data of the session**

- **Why enable_shared_from_this?**

```
struct Session
    : std::enable_shared_from_this<Session> {
    explicit Session(asio::ip::tcp::socket socket);
    void start() {
        read();
    }

private:
    void read();
    void write(std::string data);

    asio::streambuf buffer{};
    std::istream input{&buffer};
    asio::ip::tcp::socket socket;
};
```

- The session object would die at the end of the accept handler

- Thus it needs to be allocated on the heap

```
//In the accept handler
if (!ec) {
    auto session = std::make_shared<Session>(std::move(peer));
    session->start();
}
```

- The handlers need to keep the object alive

```
//In the accept handler
void Session::read() {
    auto handler = [self = shared_from_this()](error_code ec, size_t length) {
        //...
    };
}
```

```
auto session = std::make_shared<Session>(std::move(peer));  
session->start();
```

```
void Session::read() {  
    auto readCompletionHandler = [self = shared_from_this()]
```

```
void Session::write(std::string input) {  
    auto data = std::make_shared<std::string>(input);  
    auto writeCompletionHandler = [self = shared_from_this(), data]
```

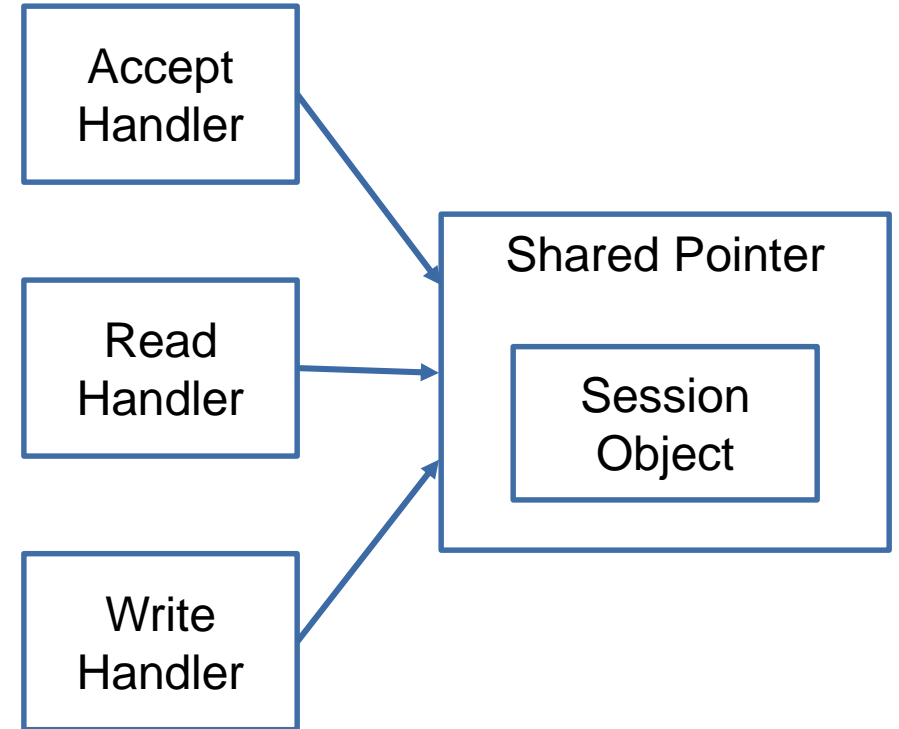
Accept
Handler

Read
Handler

Write
Handler

Shared Pointer

Session
Object



```
void WebServer::accept() {  
    auto acceptHandler = [this] (asio::error_code ec, tcp::socket peer) {  
        if (!ec) {  
            auto session = std::make_shared<Session>(std::move(peer));  
            session->start();  
        }  
    };  
    acceptor.async_accept(acceptHandler);  
}
```

```
void WebServer::accept() {  
    auto acceptHandler = [this] (asio::error_code ec, tcp::socket peer) {  
        if (!ec) {  
            auto session = std::make_shared<Session>(std::move(peer));  
            session->start();  
        }  
    };  
    acceptor.async_accept(acceptHandler);  
}
```

- **accept() should very likely be called at the end of the handler**
- **Otherwise only a single connection from a client will be possible**


```
void Session::read() {
    auto readCompletionHandler = [this] (asio::error_code ec, std::size_t length) {
        if (ec) {
            //error handling
        }
        int number{};
        if (input >> number) {
            input.ignore(std::numeric_limits<std::streamsize>::max(), '\n');
            write(createReply(number));
        }
    };
    asio::async_read_until(socket, buffer, '\n', readCompletionHandler);
}
```

```
void Session::read() {
    auto readCompletionHandler = [self = shared_from_this()] (asio::error_code ec, std::size_t length) {
        if (ec) {
            //error handling
        }
        int number{};
        if (self->input >> number) {
            self->input.ignore(std::numeric_limits<std::streamsize>::max(), '\n');
            self->write(self->createReply(number));
        }
    };
    asio::async_read_until(socket, buffer, '\n', readCompletionHandler);
}
```


- It is likely that the shared pointer to the this object needs to be captured unless the session is stored somewhere else
- It needs to live longer than the operations performed on it

Multi-Threading



- **You might find a lot of legacy C++ code using “pthread” (aka POSIX-Threads) API**
 - There is no portability guarantee
 - Your compiler must know about your use of “pthread” to generate safe code
- **C++11 removed the need to rely on POSIX-Threads**
 - Not 100% functionality equivalent for all tasks, so some implementations still use POSIX (or even Microsoft) Threads
 - BUT, it is guaranteed to be portable across platforms and compilers
 - BECAUSE
C++11 and later define a Memory Model for the execution with concurrent execution agents (aka threads)
- **The slides distinguish `std::thread` (C++ class) and `thread` (OS execution agent)**

API of `std::thread`



```
class std::thread
```

```
int main() {  
    std::thread greeter {  
        [] { std::cout << "Hello, I'm thread!" << std::endl; }  
    };  
    greeter.join();  
}
```

- A new thread is created and started automatically
- Creates a new execution context (thread)
- join() waits for the thread to finish

- Besides lambdas also functions or functor objects can be executed in a thread
- Calls the given "function" in that thread
- Return value of the function is ignored
- Threads are default-constructible and moveable (construction and assignment)

```
struct Functor {  
    void operator()() const {  
        std::cout << "Functor" << std::endl;  
    }  
};  
  
void function() {  
    std::cout << "Function" << std::endl;  
}  
  
int main() {  
    std::thread functionThread{function};  
    std::thread functorThread{Functor{}};  
  
    functorThread.join();  
    functionThread.join();  
}
```

```
template<class Function, class... Args>
explicit thread(Function&& f, Args&&...args);
```

- `std::thread` constructor takes a function/functor/lambda and arguments to forward
- You should pass all arguments by value to avoid data races and dangling references (if possible)
- Capturing by reference in lambdas creates shared data as well!

```
std::size_t fibonacci(std::size_t n) {
    if (n < 2) {
        return n;
    }
    return fibonacci(n - 1) + fibonacci(n - 2);
}

void printFib(std::size_t n) {
    auto fib = fibonacci(n);
    std::cout << "fib(" << n << ") is "
              << fib << '\n';
}

int main() {
    std::thread function { printFib, 46 };
    std::cout << "waiting..." << std::endl;
    function.join();
}
```


- **Guess what happens**

```
int main() {  
    std::thread lambda {  
        [] {std::cout << "Lambda" << std::endl; }  
    };  
    std::cout << "Main" << std::endl;  
}
```




Main

Lambda




terminate called without an active exception

- Before the `std::thread` object is destroyed you must `join()` or `detach()` the thread

```
int main() {  
    std::thread worker { doWork };  
    worker.join();  
}
```



```
int main() {  
    std::thread worker { doWork };  
    worker.detach();  
}
```



- **The destructor doesn't call `join()` nor `detach()`**
 - If it called `join()` the program might hang when leaving the scope (possibly unexpected due to an exception)
 - If it called `detach()` the thread continues with possibly deallocated resources (references to local variables)
- **If an unjoined and undetached thread object is destroyed `std::terminate()` will be called**
- **When using `.detach()` beware of the lifetime of objects referred from the detached thread's function, global or passed, or even local ones.**
 - When the `main()` thread ends globals like `std::cout` will be destroyed

```
void startThread() {  
    using namespace std::chrono_literals;  
    std::string local{"local"};  
    std::thread t{[&] {  
        std::this_thread::sleep_for(1s);  
        std::cout << local << std::endl;  
    }};  
    t.detach();  
}  
  
int main() {  
    using namespace std::chrono_literals;  
    startThread();  
    std::this_thread::sleep_for(2s);  
}
```



- **Suggested by Anthony Williams in “C++ Concurrency in Action”**
 - Adapted version from the book
- **RAII wrapper that automatically calls `join()`**
 - Probably won't be standardized

```
struct ScopedThread {  
    explicit ScopedThread(std::thread && t)  
        : the_thread{std::move(t)} {  
        if (!the_thread.joinable())  
            throw std::logic_error { "no thread" };  
    }  
    ~ScopedThread() {  
        the_thread.join();  
    }  
private:  
    std::thread the_thread;  
};  
  
int main() {  
    ScopedThread t { std::thread {  
        [] {std::cout << "Hello Thread"<< std::endl;}  
    }};  
    std::cout << "Hello Main" << std::endl;  
}
```

- Using global streams does not create data races, but sequencing of characters could be mixed

```
#include <thread>
#include <iostream>

int main() {
    using std::cout;
    using std::endl;

    std::thread t {[] {cout << "Hello Thread" << endl;}};
    cout << "Hello Main" << endl;
    t.join();
}
```

?

Hello Main
Hello Thread

HHeellllloo MTahirnead

<Arbitrary Combination>

- **namespace `std::this_thread` provides some helper functions**
- **`get_id()`, also available as member function**
 - An id of the underlying OS thread
 - Distinguishes one thread from all others and can be used as a key in a map of threads
- **`sleep_for(duration)`**
 - Suspends thread for a duration
- **`sleep_until(time_point)`**
- **`yield()`**
 - Allows OS to schedule another thread
- **NB: Timing doesn't guarantee sequence!**

```
int main() {
    using std::cout;
    using std::endl;
    using namespace std::chrono_literals;

    std::thread t { [] {
        std::this_thread::yield();
        cout << "Hello ID: "
              << std::this_thread::get_id()
              << endl;
        std::this_thread::sleep_for(10ms);
    }};
    cout << "main() ID: "
          << std::this_thread::get_id()
          << endl;
    cout << "t.get_id(): "
          << t.get_id()
          << endl;
    t.join();
}
```

```
void calcAsync() {  
    std::thread t{longRunningAction};  
    doSomethingElse();  
    t.join();  
}
```

```
void countAsync(std::string_view input) {  
    std::thread t{[&] {  
        countAs(input);  
    }};  
    t.detach();  
}
```

```
void calcAsync() {  
    std::thread t{longRunningAction};  
    doSomethingElse();  
    t.join();  
}
```

Depends

If doSomethingElse() does not throw an exception the code is correct.
If an exception is thrown you are doomed!

```
void countAsync(std::string_view input) {  
    std::thread t{[&] {  
        countAs(input);  
    }};  
    t.detach();  
}
```

Incorrect

The value parameter is captured by reference. It runs out of scope after the function execution. Furthermore, string_view is a reference wrapper itself. The string it is referring to might be destroyed as well.

- **Until the C++ standard specifies the networking TS (>C++20) you have to rely on external implementations (like ASIO)**
- **ASIO is best used to implement asynchronous operations (available in boost or standalone)**
- **The asynchronous programming model has special pitfalls in non-garbage collected environments**
- **The C++ standard features API for using threads**