

Department I - C Plus Plus

Modern and Lucid C++ Advanced for Professional Programmers

Week 14 – WebAssembly

Thomas Corbat / Felix Morgner
Rapperswil / St. Gallen, 06.06.2024
FS2024



OST

Ostschweizer
Fachhochschule

- **Recap Week 13**
- **Introduction to WebAssembly**

- **Participants should ...**
 - ... know the basics of WebAssembly and WebAssembly Text
 - ... be able to compile a C++ application to WebAssembly
 - ... be able to run the application in a browser

Recap Week 13



- **Five libraries**
 - All depending on a common infrastructure library
- **Two executables**
 - Depend on some or all of the libraries
- **Two target-platforms**
 - Linux on x86_64 and armv7
 - OS X
- **Code will change owners**
- **4 months time-frame**

- **Write a script that...**

- Compiles each source file
- Links all object files together
- Repeats that for every target

- **DON'T! Because...**

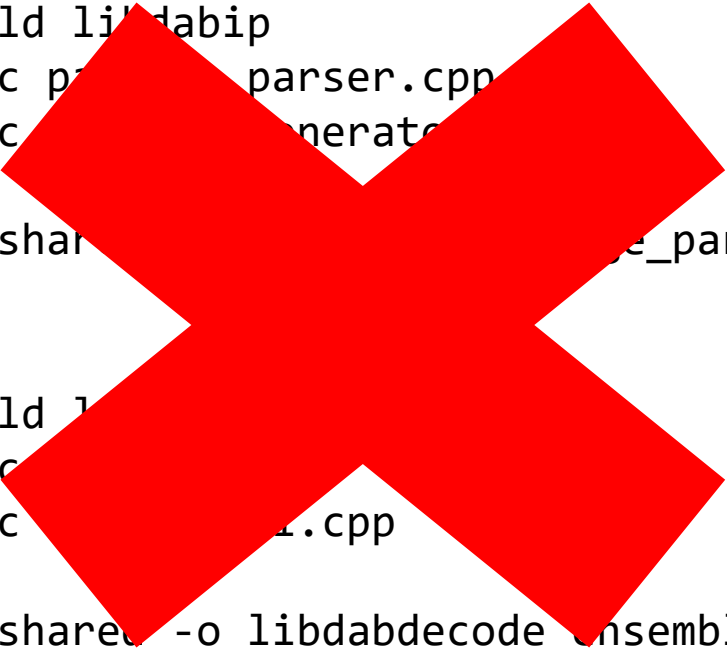
- ... every source file get built every time!
- ... the commands tend to be platform specific
- ... build order must be managed manually
- ... scripts tend to become messy over time

```
#!/bin/bash

# Build libdabip
gcc -c parser.cpp
gcc -c generator.cpp
...
gcc -shared -o libdabip.o parser.o ...

# Build 1
gcc -c 1.cpp
gcc -c 2.cpp
...
gcc -shared -o libdabdecode ensemble.o ...

# Build <you get the idea>
```



- **Make-style Build Tools**

- Run build scripts
- Produce your final products
- Often verbose
- Use a language agnostic configuration language

- **Build Script Generators**

- Generate configurations for Make-style Build Systems or Build Scripts
- Configuration independent of actual build tool
- Advanced features (download dependencies, etc.)

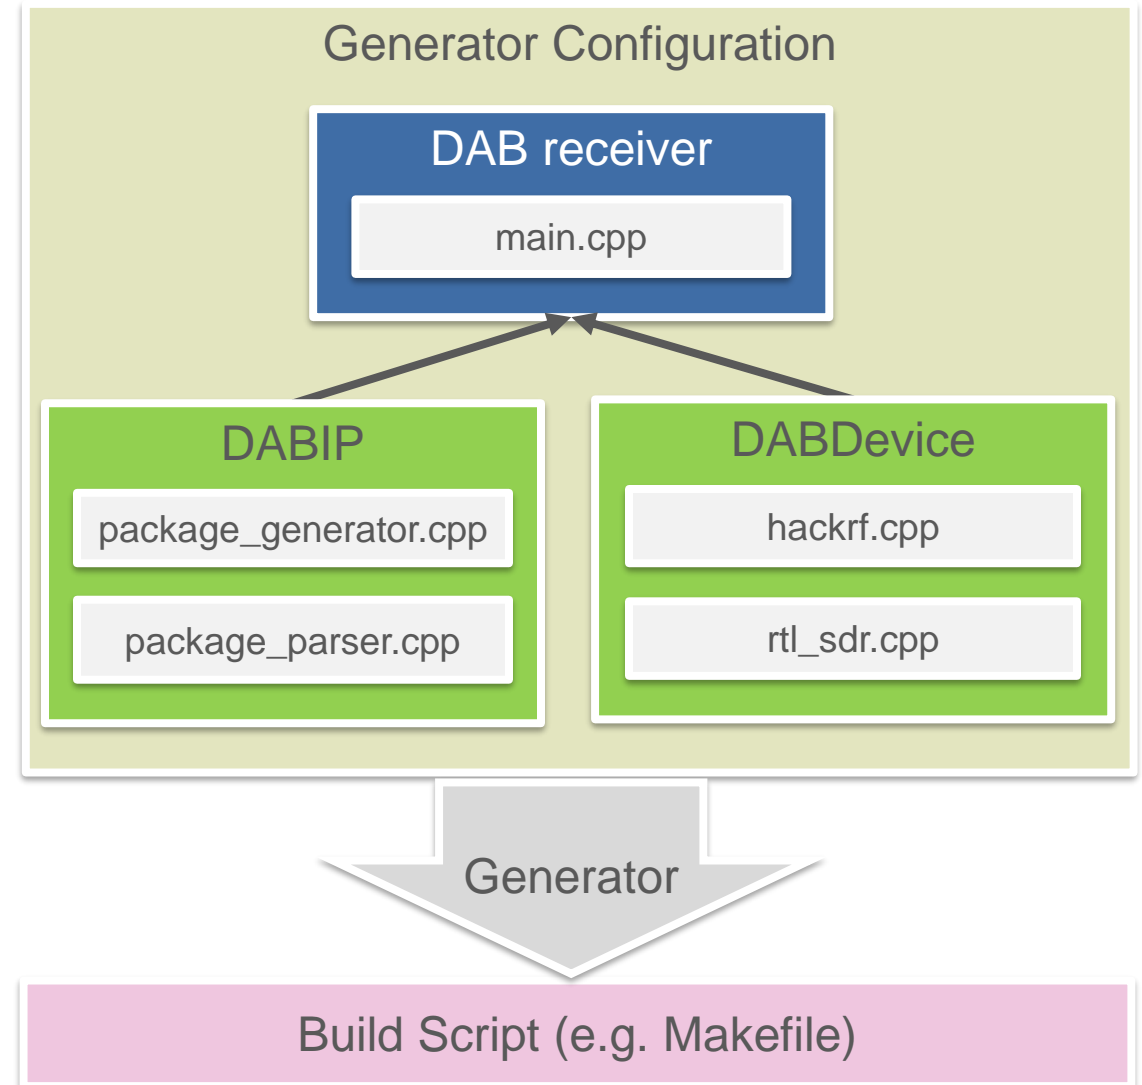
- **Idea: Take a step back**

- Define what we want to achieve, not how to do it
- Work on a higher level
- Let the create the actual build configurations

- **Platform independent build specification**

- **Tool Independent**

- Often can generate IDE projects
- Support multiple build tools



- **CMake includes CTest**

- Enable CTest using `enable_testing()`
- Create a “Test Runner” executable
 - Make sure to include your suite sources!

- Configure build environment:

```
$ cmake ..
```

- Build the project:

```
$ cmake --build .
```

- Run ctest

```
$ ctest --output-on-failure
```

```
cmake_minimum_required(VERSION "3.12.0")

project("answer" LANGUAGES CXX)

enable_testing()

add_library("${PROJECT_NAME}"
  "answer.cpp"
)

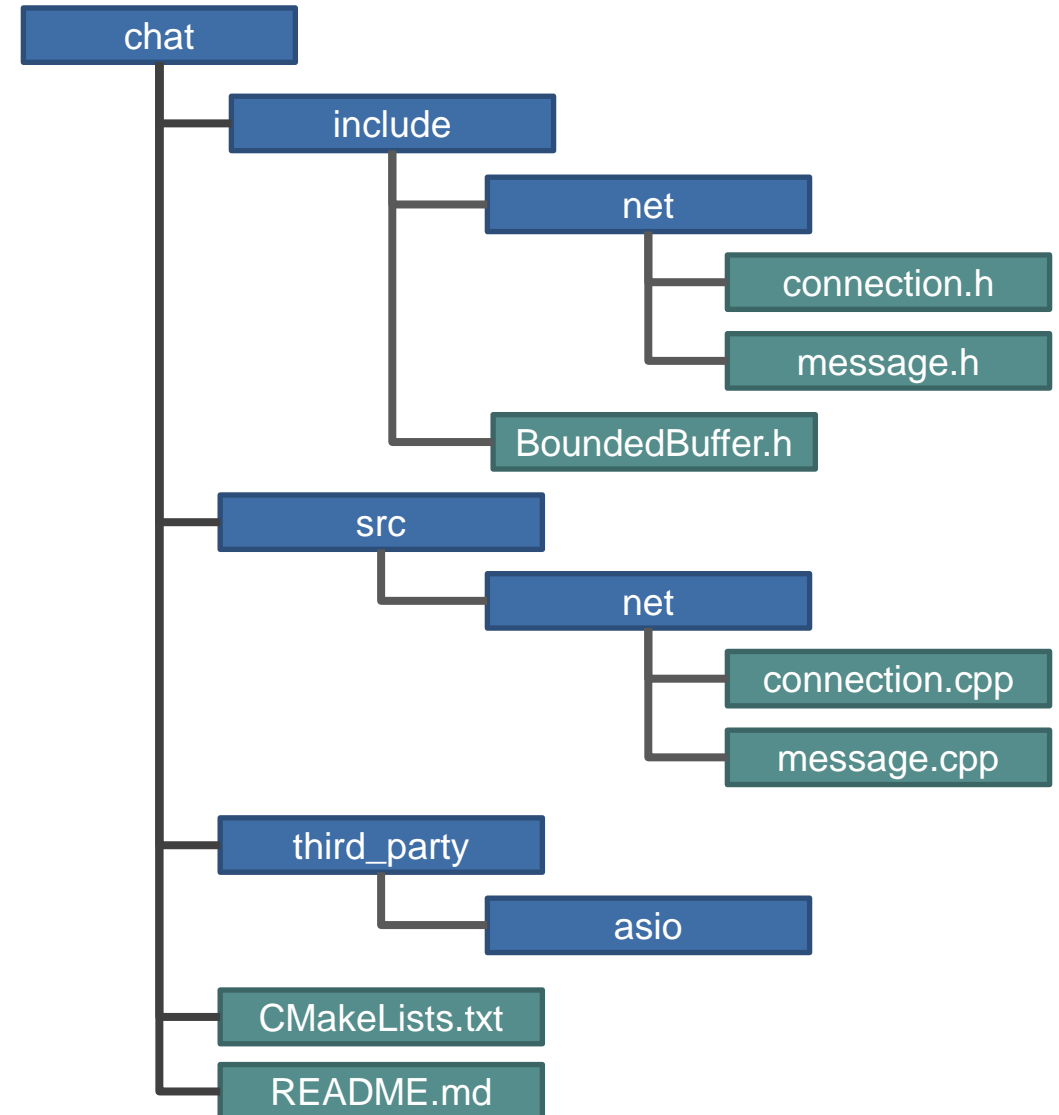
add_executable("test_runner"
  "Test.cpp"
)

target_link_libraries("test_runner" PRIVATE
  "answer"
)

target_include_directories("test_runner" SYSTEM PRIVATE
  "cute"
)

add_test("tests" "test_runner")
```

- **Headers live in the “include” folder**
 - Add subfolders for separate subsystems if needed
- **Implementation files live in the “src” folder**
 - Make sure that subfolder layout matches the “include” folder (**consistency**)
- **Put third-party projects/sources in a “third_party” or “lib” folder**
- **Test resource live in the “test” folder**
 - The test folder will have src, include, and third_party subfolders if required
- **Build configuration files should live in the root of your project**



- **Libraries may benefit from a slightly different layout**

- You will need to ship your headers
- Your headers might have very generic names

- **Idea: Introduce another nesting level for your headers**

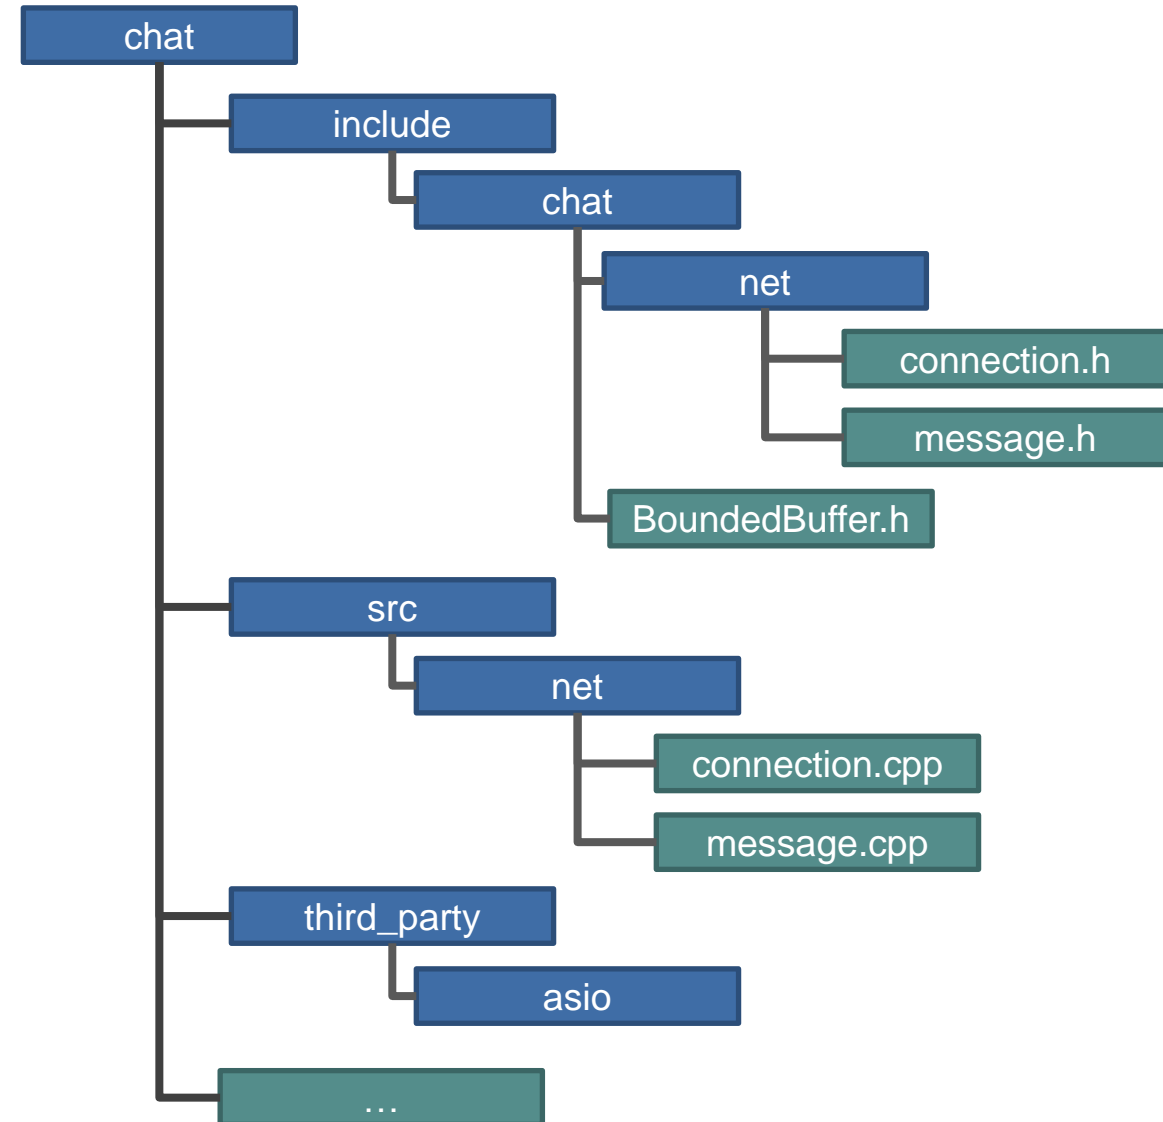
- Use the name of your project

```
#include "net/message.h"
```

... becomes ...

```
#include "chat/net/message.h"
```

- Helps avoid filename clashes



WebAssembly Basics

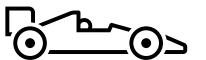
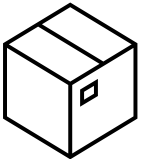
Based on The Art of WebAssembly



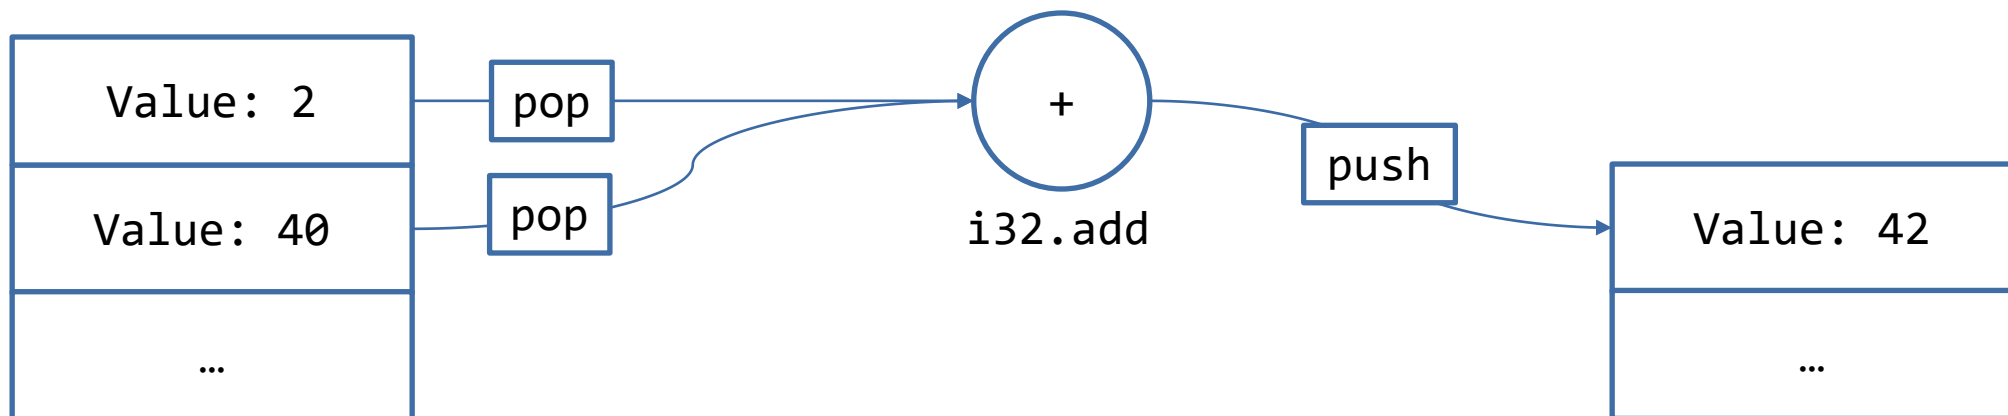
- **Can be run in modern web browsers**
- **No plug-in required (opposed to Java in browsers)**
- **Low-level language with compact representation**
- **High performance (number crunching, rendering, etc.)**
- **Enables cross-compilation of languages like C++ (Emscripten) and Rust (wasm-pack) to the web**
- **Integrates well with JavaScript**



- **Instruction Set Architecture for a stack machine (similar to JVM)**
- **Targets a virtual machine, which can be implemented for various physical machines**
- **Small binary size (for fast download) and portable**
- **Significantly faster than corresponding JavaScript code (in the browser)**
 - Compiled and optimized upfront -> binary format
 - JavaScript is text that needs to be parsed, interpreted, JIT compiled and optimized on the fly
- **Less versatile for the programmer as its features is limited compared to JavaScript**
 - Requires another high-level language as source for reasonable applications
 - Best suited for high-performance use cases



- **No registers (as in register machines)**
- **Operations are performed on the topmost elements of a stack**
 - Pop elements
 - Apply operation
 - Push result



- **WebAssembly is distributed in binary format**
- **WebAssembly Text is a mnemonic form of WebAssembly (its disassembly)**
 - (Super)human-readable
- **Instructions that work on the stack (push, pop, both)**
 - Instruction format <Type>.<Operation>
 - Example: i32.add
- **Declarations with keywords**
 - Module, global, func, etc.

```
(module
  (global $a_val (mut i32) (i32.const 1))
  (global $b_val (mut i32) (i32.const 2))
  (global $c_val (mut i32) (i32.const 0))
  (func $main (export "main")
    global.get $a_val
    global.get $b_val
    i32.add
    global.set $c_val
  )
)
```


- **Alternative syntax for structuring the code**

```
(module
  (global $a_val (mut i32) (i32.const 1))
  (global $b_val (mut i32) (i32.const 2))
  (global $c_val (mut i32) (i32.const 0))
  (func $main (export "main")
    (global.set $c_val
      (i32.add (global.get $a_val) (global.get $b_val))
    )
  )
)
```

S-Expression

⇔

```
global.get $a_val
global.get $b_val
i32.add
global.set $c_val
```

Linear Syntax

- **Introduced by func**

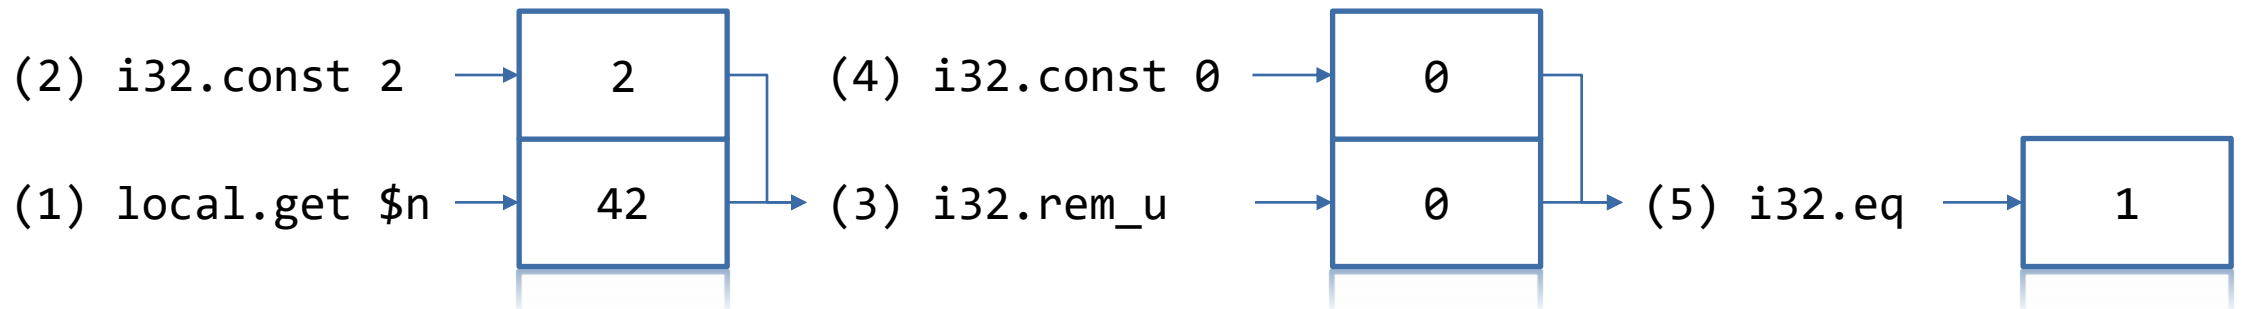
- Can be declared as “export”

- **Name after \$**

- **Param(s) with type**

- **Result with type**

```
(module
  (func $even_check (param $n i32) (result i32)
    local.get $n
    i32.const 2
    i32.rem_u
    i32.const 0
    i32.eq
  )
)
```



- **Integers**

- i32 – 32 bit
- i64 – 64 bit
- 2s Complement

- **Floating point**

- f32 – 32 bit
- f64 – 64 bit
- IEEE 754 representation

- **On JavaScript side the numbers will be mapped to 64-bit float**

- **.eq** – Test for equality
- **.ne** – Test not equal
- **.lt[_s|_u]** – Test for less than
- **.le[_s|_u]** – Test for less or equal
- **.gt[_s|_u]** – Test for greater than
- **.ge[_s|_u]** – Test for greater or equal
- **.and** – Bitwise AND
- **.or** – Bitwise OR
- **.xor** – Bitwise XOR
- **.eqz** – Float equals zero

- **_s** or **_u** suffix for signed and unsigned comparison
- **Consumes the top two elements on the stack and pushes the result**
- **Example**

```
local.get $x  
local.get $y  
i32.gt_s ;; pushes 1 on the stack if $x > $y
```

- **.add – Plus operation**
- **.sub – Minus operation**
- **.mul – Multiply operation**
- **.div[_s|_u] – Division operation**
- **.rem[_s|_u] – Remainder operation**

- **Consumes the top two elements on the stack and pushes the result**

- **Example**

```
local.get $x  
local.get $y  
i32.mul ;; pushes $x * $y on the stack
```

- **.and** – Bitwise AND
- **.or** – Bitwise OR
- **.xor** – Bitwise XOR
- **.eqz** – Float equals zero

- **Consumes the top two elements on the stack and pushes the result**

- **Example**

```
local.get $x  
local.get $y  
i32.xor ;; pushes  $\$x \wedge \$y$  to the stack
```

- **if...else**

```
local.get $condition_i32
if
    ;; $condition_i32 not 0
    nop
else
    ;; $condition_i32 is 0
    nop
end
```

- **loop**

```
(loop $infinite_loop
    nop
    br $infinite_loop ;; jump to loop label
)
```

- **block**

```
(block $jump_to_end
    br $jump_to_end ;; jump to end of block
    nop ;; never executed
)
```

- **Every WAT application must be a module**
- **Top-level declaration**
- **;; is a line-comment**
- **(; ;) block comment (multi-line)**
 - Can be nested

```
(module  
  ;; code  
)
```


- **Allocate one page of 64KB memory**

(memory 1)

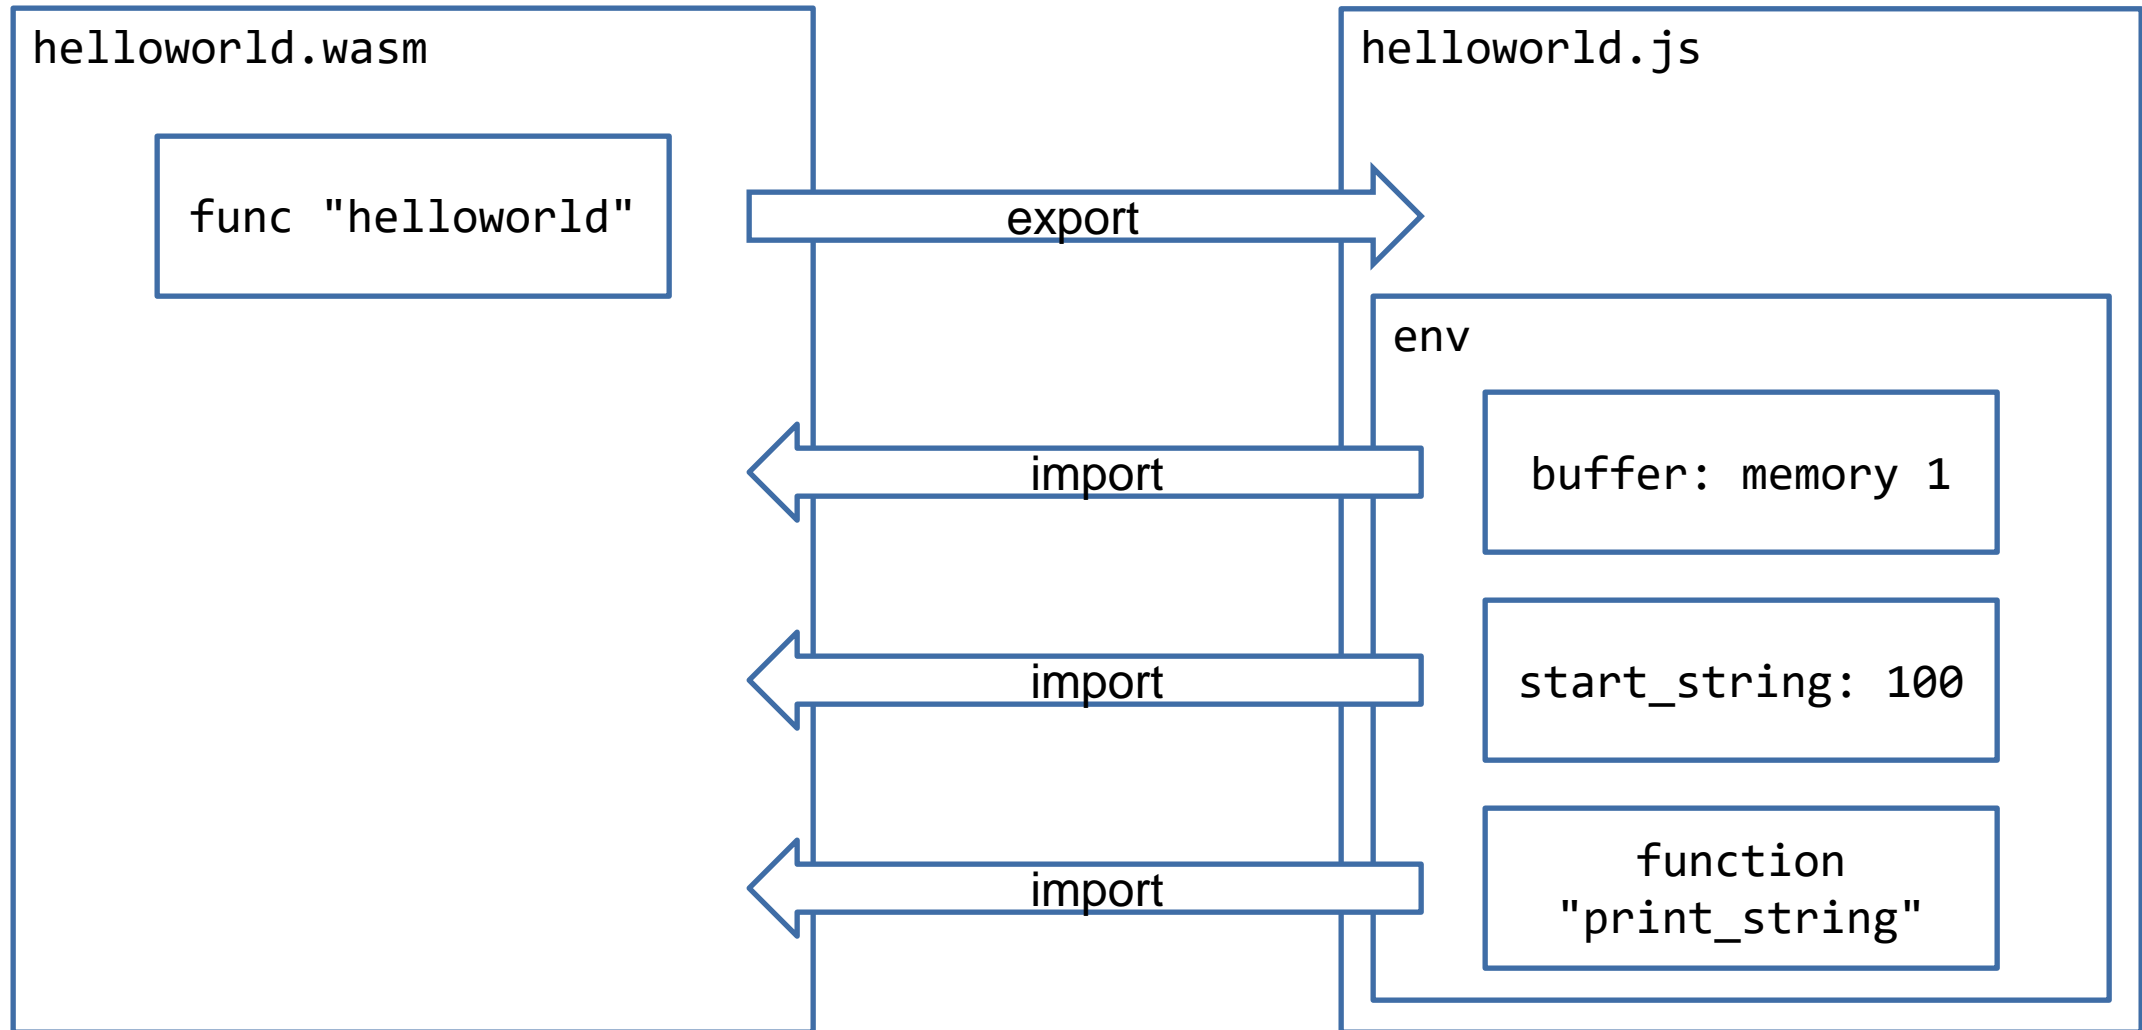
- **Maximum number of pages in an application is 65'536, i.e., 4GB memory**
 - Technical limitation as 32-bit addressing is used internally

```
(module
  (memory 1)
  (global $pointer i32 (i32.const 128))
  (func $init
    (i32.store
      (global.get $pointer) ;; store at address $pointer
      (i32.const 99) ;; value stored
    )
  )
  ;;...
  (start $init)
)
```

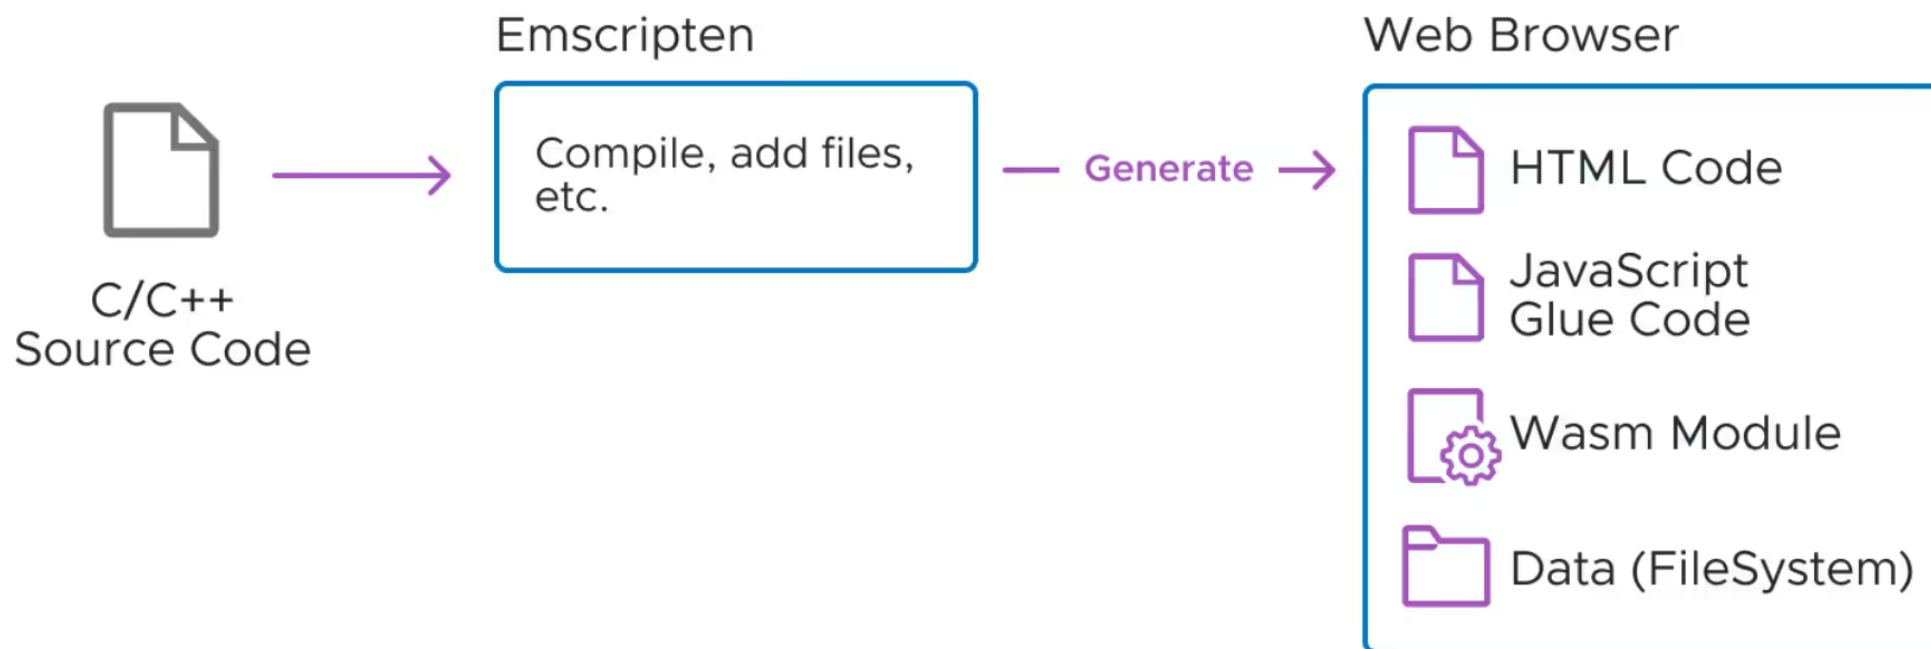
```
(module
  (import "env" "print_string" (func $print_string (param i32)))
  (import "env" "buffer" (memory 1))
  (global $start_string (import "env" "start_string") i32)
  (global $string_len i32 (i32.const 12))
  (data (global.get $start_string) "hello world!")
  (func (export "helloworld")
    (call $print_string (global.get $string_len))
  )
)
```

- **Compiled with wat2wasm**

```
const fs = require('fs');
const bytes = fs.readFileSync(__dirname + '/helloworld.wasm');
let hello_world = null; // function will be set later
let start_string_index = 100; // linear memory location of string
let memory = new WebAssembly.Memory({ initial: 1 }); // linear memory
let importObject = {
  env: {
    buffer: memory,
    start_string: start_string_index,
    print_string: function (str_len) {
      const bytes = new Uint8Array (memory.buffer, start_string_index, str_len);
      const log_string = new TextDecoder('utf8').decode(bytes);
      console.log (log_string);
    }
  }
};
( async () => {
  let obj = await
    WebAssembly.instantiate(new Uint8Array (bytes), importObject);
  ({helloworld: hello_world} = obj.instance.exports);
  hello_world();
})();
```



- Emscripten: Toolchain for LLVM-based languages to WebAssembly
- Compiler Frontend (emcc)
- <https://emscripten.org/>



- **C++ Standard Library**

- **Containers**
- **Algorithms**
- **Memory Management** (`unique_ptr` etc.)

- **Exceptions**

- **Default disabled** -> terminate!
- **Optional via JavaScript Exceptions**
- **WebAssembly Exception Proposal**

- **POSIX Networking API via WebSockets**

- **Graphics APIs**

- **OpenGL / EGL translated to WebGL**
- **SDL**

- **Multithreading**

- **Support for POSIX Threads**
- **Atomic Operations**

- **File I/O via embedded file system**

- **OpenAL support via WebAudio**

```
#include <iostream>

auto main() -> int {
    std::cout << "CplA on the Web!\n";
}
```

- `em++ -std=c++20 hello.cpp`
 - Generates WASM and JS file
 - JS file provides initialization
 - Executable with NodeJS or similar
 - Useful in backend
- `em++ -std=c++20 hello.cpp -o hello.html`
 - Generates HTML shell
 - Can be loaded in browser
 - May require local webserver!
 - Useful in frontend

● Autotools

- Emscripten provides “wrappers”
- `emconfigure`
 - Wraps execution of configure scripts
 - E.g. `emconfigure ./configure`
- `emmake`
 - Wraps execution of make itself
 - E.g. `emmake make all`

● Cmake

- Emscripten provides a single “wrapper”
- `emcmake`
 - Wraps execution of `cmake`
 - Provides additional flags under the hood
 - `emcmake cmake -S . -B build`
- Can use libraries just as usual
- Can generate HTML shell by specifying suffix
 - ```
set_target_properties(tgt_name
 PROPERTIES
 SUFFIX ".html"
)
```



- **WebAssembly is designed to be as close to “native” performance as possible while still being easily portable**
  - Enables high performance execution in web browsers as well as in backend applications using NodeJS.
  - However: It is currently limited to 32-Bit applications, meaning that memory is limited to 4GB.
- **The design of WebAssembly makes it possible to use different high-level languages to write applications for the web.**
- **JavaScript and WebAssembly can interact with each other**
  - JavaScript can call WebAssembly functions, and share data with it, and vice versa.
- **Emscripten provides a toolchain enabling the use of modern C++ on the web**
  - It integrates well with existing build systems like GNU Autotools or CMake.
  - Makes it possible to ship the same code natively or via the web!